
COGS 181, Fall 2017

Neural Networks and Deep Learning

Zhuowen Tu

Department of Cognitive Science
University of California, San Diego

Syllabus

Syllabus

Lecture Time:

Tuesday and Thursday, 12:30p-1:50pm CSB 001

Lab hours:

Wednesday 4:00p-4:50p CSB 005
Friday 1:00p-1:50p CSB 005

TA:

Zeyu Chen (zec003@eng.ucsd.edu)
Long Jin (longjin@eng.ucsd.edu)

IA:

Adilijiang Ainihaer (aainihae@ucsd.edu)

IA:

Text Books:

This course will be self-contained but below are some useful books that are good to read:

Deep Learning (Ian Goodfellow and Yoshua Bengio and Aaron Courville)

Neural Networks for Machine Learning (Geoffrey Hinton)

R. Duda, P. Hart, D. Stork, "Pattern Classification", second edition, 2000.

Web Resources:

Piazza: <https://piazza.com/configure-classes/fall2017/cogs181>

Ted: ted.ucsd.edu

Podcast: <https://podcast.ucsd.edu/>

Office Hours:

Zhuowen Tu, Tuesday and Thursday 2:00-3:00pm, CSB 132

Course Description:

Recent developments in deep neural networks approaches have led to a significant boost to state-of-the-art learning systems in a wide range of domains including machine learning, robotics, perception, computer vision, artificial intelligence, speech recognition, neural computation, medical imaging, bio-informatics, computational linguistics, and social data analysis. This course will cover the basics about neural networks, as well as recent developments in deep learning including deep belief nets, convolutional neural networks, recurrent neural networks, long-short term memory, and reinforcement learning. We will study details of the deep learning architectures with a focus on learning end-to-end models for these tasks, particularly image classification. Students will learn hands-on skills to do cutting-edge research by implementing, training and debugging neural networks.

Prerequisites:

Mathematics 20F (Linear Algebra) or Mathematics 31AH (Honors Linear Algebra), and Mathematics 180A (Introduction to Probability) or ECE 109 (Engineering Probability & Statistics), and COGS 109 (Modeling and Data Analysis) or COGS 108 or CSE 11 (Introduction to Computer Science & Object-Oriented Programming: Java), or consent of instructor.

Grading policy:

Assignments: 40% (Late policy: 5% reduction for the first day past due and every 10% reduction afterwards).

Midterms: 40%

Final project: 20% (Late policy: 5% reduction for the first day past due and every 10% reduction afterwards.).

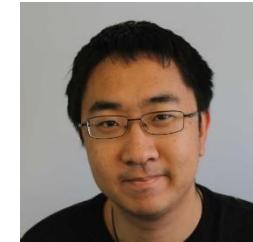
Bonus points: 3% (classroom participation and final project)

Policy on Integrity of Scholarship

TA:



Zeyu Chen



Long Jin

IA:



Adilijiang Ainihaer

Grading Policy

Assignments: 40%

Midterm exams: 40%

Final project: 20%

Bonus points: 3% (classroom participation + final project)

Course Webpage

<https://sites.google.com/site/ucsdcogs181fall2017/>

Piazza

<https://piazza.com/ucsd/fall2017/cogs181/>

Ted: ted.ucsd.edu

Class schedule

Calendar and Class Notes

Date and Lecture #	Topic	Readings			
Week 1 09/28/17 (Thursday)	Course overview and introduction Concepts about learning and neural networks	Probability theory (by Matthew Shum) Introduction to probability by C.M. Grinstead and J.L. Snell Linear algebra review A few useful things to know about machine learning (Pedro Domingos) Learning Python Why do we need machine learning? What are neural networks? (G. Hinton)			
Week 2 10/03/17 (Tuesday)	Data Classification Discriminative vs Generative learning	Deep Learning: Linear Algebra Deep Learning: Machine Learning Basics			
Week 2 10/05/17 (Thursday)	Decision boundaries Introduction to neurons and neural networks	Deep Learning: Numeric Optimization			
Week 3 10/10/17 (Tuesday)	Error metrics and optimization Least square estimation Convexity	Precision and recall Receive operating characteristic Least square estimation Convex function Matrix calculus			
Week 3 10/12/17 (Thursday)	Gradient descent and logistic regression Perceptron: training and testing	Gradient descent Logistic regression Perceptron			
Week 4 10/17/17 (Tuesday)	Hopfield neural networks	Hopfield Neural Network J. J. Hopfield and D. W. Tank, "Neural" computation of decisions in optimization problems, Biological Cybernetics 1985.			
Week 4 10/19/17 (Thursday)	Midterm 1 Hopfield neural networks				
			Week 5 10/24/17 (Tuesday) Week 5 10/26/17 (Thursday) Week 6 10/31/17 (Tuesday) Week 6 11/02/17 (Thursday) Week 7 11/07/17 (Tuesday) Week 7 11/09/17 (Thursday) Week 8 11/14/17 (Tuesday) Week 8 11/16/17 (Thursday) Week 9 11/21/17 (Tuesday) Week 9 11/23/17 (Thursday) Week 10 12/05/17 (Tuesday) Week 10 12/07/17 (Thursday)	Feedforward networks Feedforward networks Convolutional neural networks Convolutional neural networks: learning strategies, applications Convolutional neural networks Modern CNN architectures Deep belief nets Midterm 2 Auto-encoder Language modeling and Word2Vec Recurrent neural networks No class Long-short memory Deep reinforcement learning	Feedforwad neural networks Backpropagation Convolution Circuit illustration of backpropagation AlexNet GoogLeNet VGG ResNet Boltzman machine Restricted Boltzman manchine Word2Vec Recurrent neural network Long short term memory Gated recurrent unit

Background materials to learn and prepare

Linear algebra review:

<http://cs229.stanford.edu/section/cs229-linalg.pdf>

Probability theory

<http://people.hss.caltech.edu/~mshum/stats/lect1.pdf>

Python/IPython

<http://ipython.org/>

<http://cs231n.github.io/python-numpy-tutorial/>

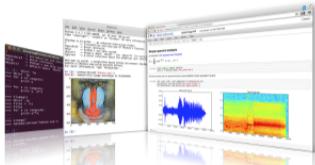
IPython Notebook

IP[y]: IPython Interactive Computing

[Install](#) · [Documentation](#) · [Project](#) · [Jupyter](#) · [News](#) · [Cite](#) · [Donate](#) · [Books](#)

IPython provides a rich architecture for interactive computing with:

- A powerful interactive shell.
- A kernel for [Jupyter](#).
- Support for interactive data visualization and use of [GUI toolkits](#).
- Flexible, [embeddable](#) interpreters to load into your own projects.
- Easy to use, high performance tools for [parallel computing](#).



To get started with IPython in the Jupyter Notebook, see our [official example collection](#). Our [notebook gallery](#) is an excellent way to see the many things you can do with IPython while learning about a variety of topics, from basic programming to advanced statistics or quantum mechanics.

To learn more about IPython, you can download our [talks and presentations](#), or read our [extensive documentation](#). IPython is open source (BSD license), and is used by a range of [other projects](#); add your project to that list if it uses IPython as a library, and please don't forget to [cite the project](#).

IPython supports Python 2.7 and 3.3 or newer. Our older 1.x series supports Python 2.6 and 3.2.

Jupyter and the future of IPython

IPython is a growing project, with increasingly language-agnostic components. IPython 3.x was the last monolithic release of IPython, containing the notebook server, qtconsole, etc. As of IPython 4.0, the language-agnostic parts of the project: the notebook format, message protocol, qtconsole, notebook web application, etc. have moved to new projects under the name [Jupyter](#). IPython itself is focused on interactive Python, part of which is providing a Python kernel for Jupyter.

Announcements

- [IPython 6.1](#), and [IPython 5.4](#): These two releases were made on May 31st 2017.
- [IPython 6.0](#): This release, the first to require Python 3, integrates the Jedi library for completion. See the [release notes](#) for more information about what's new.
- [JupyterCon 2017](#): The first Jupyter Community Conference will take place in New York City on August 23-25 2017, along with a satellite training program on August 22-23. The Project Jupyter team has partnered with O'Reilly Media for this event; for more details, including submitting a talk, see the [JupyterCon website](#).
- [IPython 5.0](#): The release of IPython 5.0 brings a major revision of the terminal interface, including syntax highlighting as you type and better multiline editing, thanks to the [prompt_toolkit](#) library. See the [release notes](#) for more about the new features.
- [Book](#): Cyrille Rossant has published the second edition of the IPython minibook: [Learning IPython for Interactive Computing and Data Visualization](#), for which [Damian Avila](#) was a technical reviewer. We thank Packt Publishing for donating a portion of the proceeds from this book to support IPython's development.
- [O'Reilly Book](#): [Mining the Social Web](#) is an open source data science project and [book](#) that features nearly 130 examples with IPython Notebook and a Vagrant-powered virtual machine environment. You can preview all of the example notebooks from its GitHub [repository](#) on IPython's Notebook Viewer [here](#).

[More news...](#)

IP[y]: IPython Interactive Computing

[Install](#) · [Documentation](#) · [Project](#) · [Jupyter](#) · [News](#) · [Cite](#) · [Donate](#) · [Books](#)

Installing IPython

There are multiple ways of installing IPython. This page contains simplified installation instructions that should work for most users. Our official documentation contains [more detailed instructions](#) for manual installation targeted at advanced users and developers.

If you are looking for installation documentation for the notebook and/or qtconsole, those are now part of [Jupyter](#).

I already have Python

If you already have Python installed and are familiar with installing packages, you can get IPython with pip:

```
pip install ipython
```

I am getting started with Python

For new users who want to install a full Python environment for scientific computing and data science, we suggest installing the Anaconda or Canopy Python distributions, which provide Python, IPython and all of its dependences as well as a complete set of open source packages for scientific computing and data science.

1. Download and install Continuum's [Anaconda](#) or the free edition of Enthought's [Canopy](#).
2. Update IPython to the current version using the Terminal:

Anaconda:

```
conda update conda
conda update ipython
```

Enthought Canopy:

```
enpkg ipython
```

Downloads

You can manually download IPython from [GitHub](#) or [PyPI](#). To install one of these versions, unpack it and run the following from the top-level source directory using the Terminal:

```
pip install .
```

Stanford CS231n: <http://cs231n.stanford.edu/>

CS231n: Convolutional Neural Networks for Visual Recognition
Spring 2017



*This network is running live in your browser

Course Description

Computer Vision has become ubiquitous in our society, with applications in search, image understanding, apps, mapping, medicine, drones, and self-driving cars. Core to many of these applications are visual recognition tasks such as image classification, localization and detection. Recent developments in neural network (aka "deep learning") approaches have greatly advanced the performance of these state-of-the-art visual recognition systems. This course is a deep dive into details of the deep learning architectures with a focus on learning end-to-end models for these tasks, particularly image classification. During the 10-week course, students will learn to implement, train and debug their own neural networks and gain a detailed understanding of cutting-edge research in computer vision. The final assignment will involve training a multi-million parameter convolutional neural network and applying it on the largest image classification dataset (ImageNet). We will focus on teaching how to set up the problem of image recognition, the learning algorithms (e.g. backpropagation), practical engineering tricks for training and fine-tuning the networks and guide the students through hands-on assignments and a final course project. Much of the background and materials of this course will be drawn from the ImageNet Challenge.

Instructors



Fei-Fei Li



Justin Johnson



Serena Yeung

Teaching Assistants



Albert Haque
(Head TA)



Rishi Bedi



Shyamal Buch



Zhao (Joe) Chen



Timnit Gebru



Agrim Gupta



De-An Huang



Russell Kaplan



Leo Keselman



Nishith Khendwala



Xingyu Liu



Shayne Longpre



Zejun (Alan) Luo



Lane McIntosh



Olivier Moindrot



Amani Peddada



Boya (Emma)
Peng



Ben Poole



Luda Zhao

Lecture videos: <https://www.youtube.com/playlist?list=PL3FW7Lu3i5JvHM8ljYj-zLfQRF3EO8sYv>

IPython tutorial to quickly get started

CS231n Convolutional Neural Networks for Visual Recognition

Python Numpy Tutorial

This tutorial was contributed by [Justin Johnson](#).

We will use the Python programming language for all assignments in this course. Python is a great general-purpose programming language on its own, but with the help of a few popular libraries (numpy, scipy, matplotlib) it becomes a powerful environment for scientific computing.

We expect that many of you will have some experience with Python and numpy; for the rest of you, this section will serve as a quick crash course both on the Python programming language and on the use of Python for scientific computing.

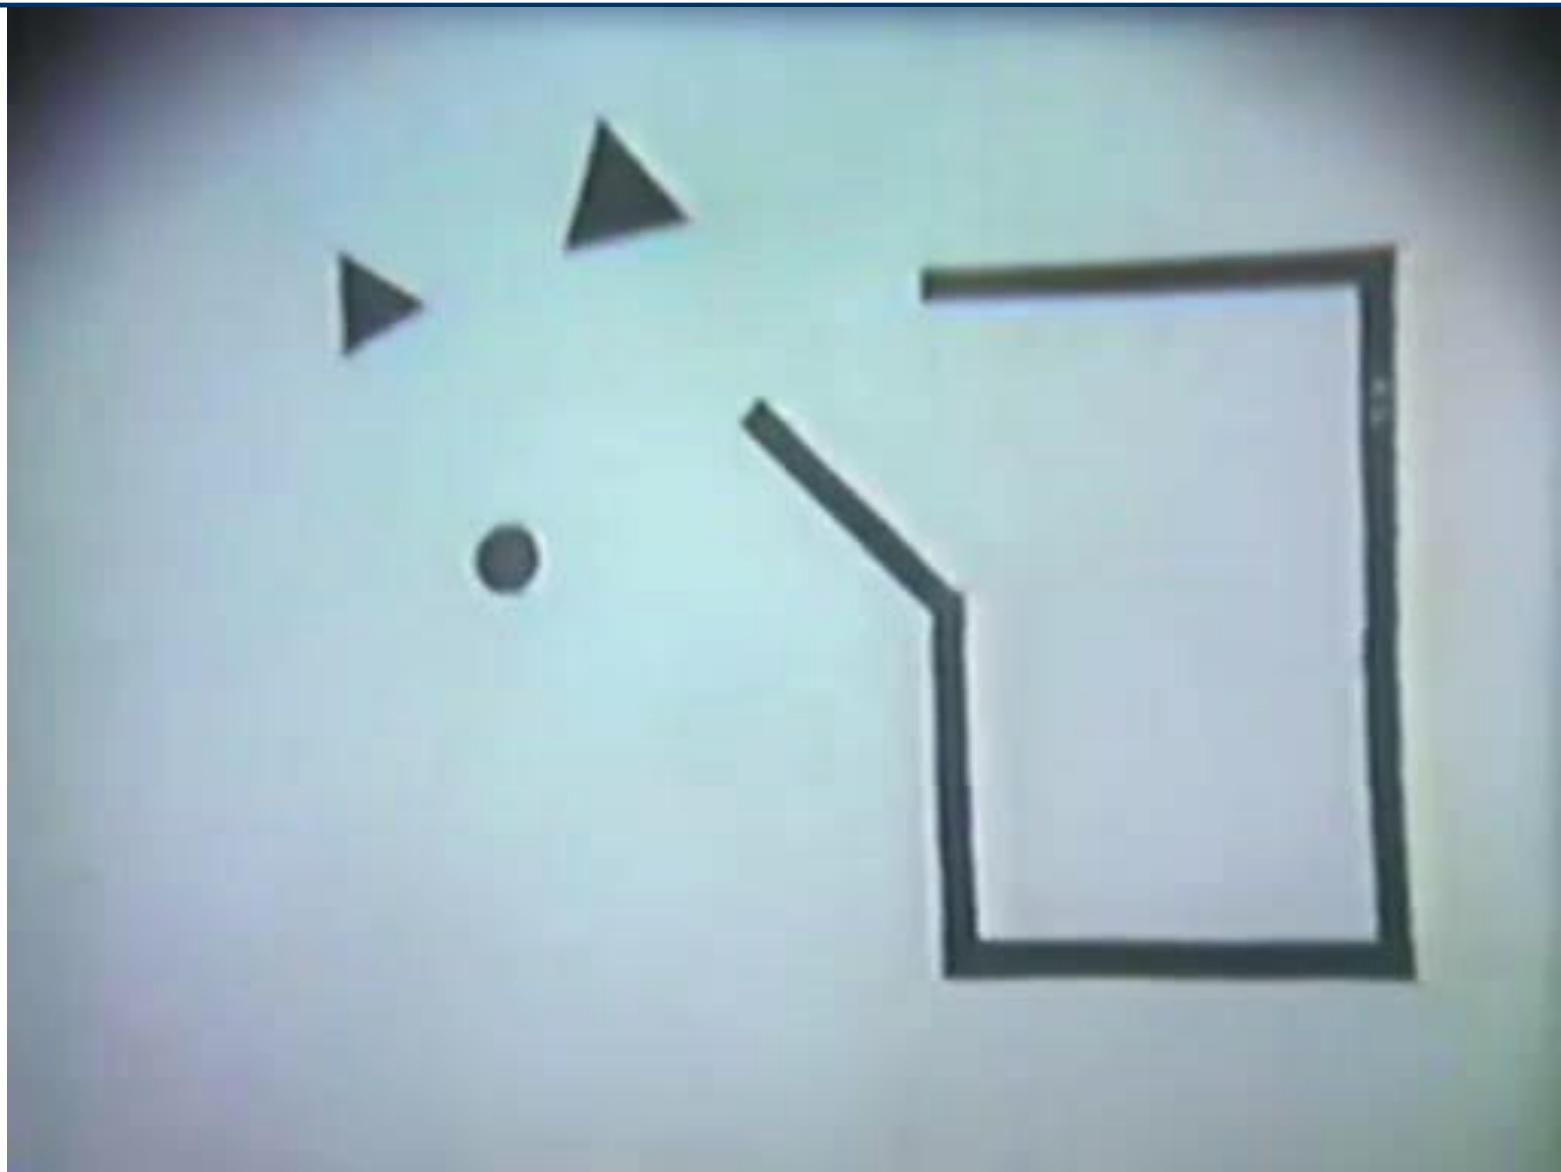
Some of you may have previous knowledge in Matlab, in which case we also recommend the [numpy for Matlab users](#) page.

You can also find an [IPython notebook version of this tutorial here](#) created by [Volodymyr Kuleshov](#) and [Isaac Caswell](#) for CS 228.

Table of contents:

- [Python](#)
 - [Basic data types](#)
 - [Containers](#)
 - [Lists](#)
 - [Dictionaries](#)
 - [Sets](#)
 - [Tuples](#)
 - [Functions](#)
 - [Classes](#)
- [Numpy](#)
 - [Arrays](#)
 - [Array indexing](#)
 - [Datatypes](#)
 - [Array math](#)
 - [Broadcasting](#)
- [SciPy](#)
 - [Image operations](#)
 - [MATLAB files](#)
 - [Distance between points](#)
- [Matplotlib](#)
 - [Plotting](#)
 - [Subplots](#)
 - [Images](#)

Why is representation so important?



Video: Apparent Behavior, Heider and Simmel, 1944

Artificial neural networks: a brief history

Rosenblatt, F. (1958). "The Perceptron: A Probabilistic Model For Information Storage And Organization In The Brain".

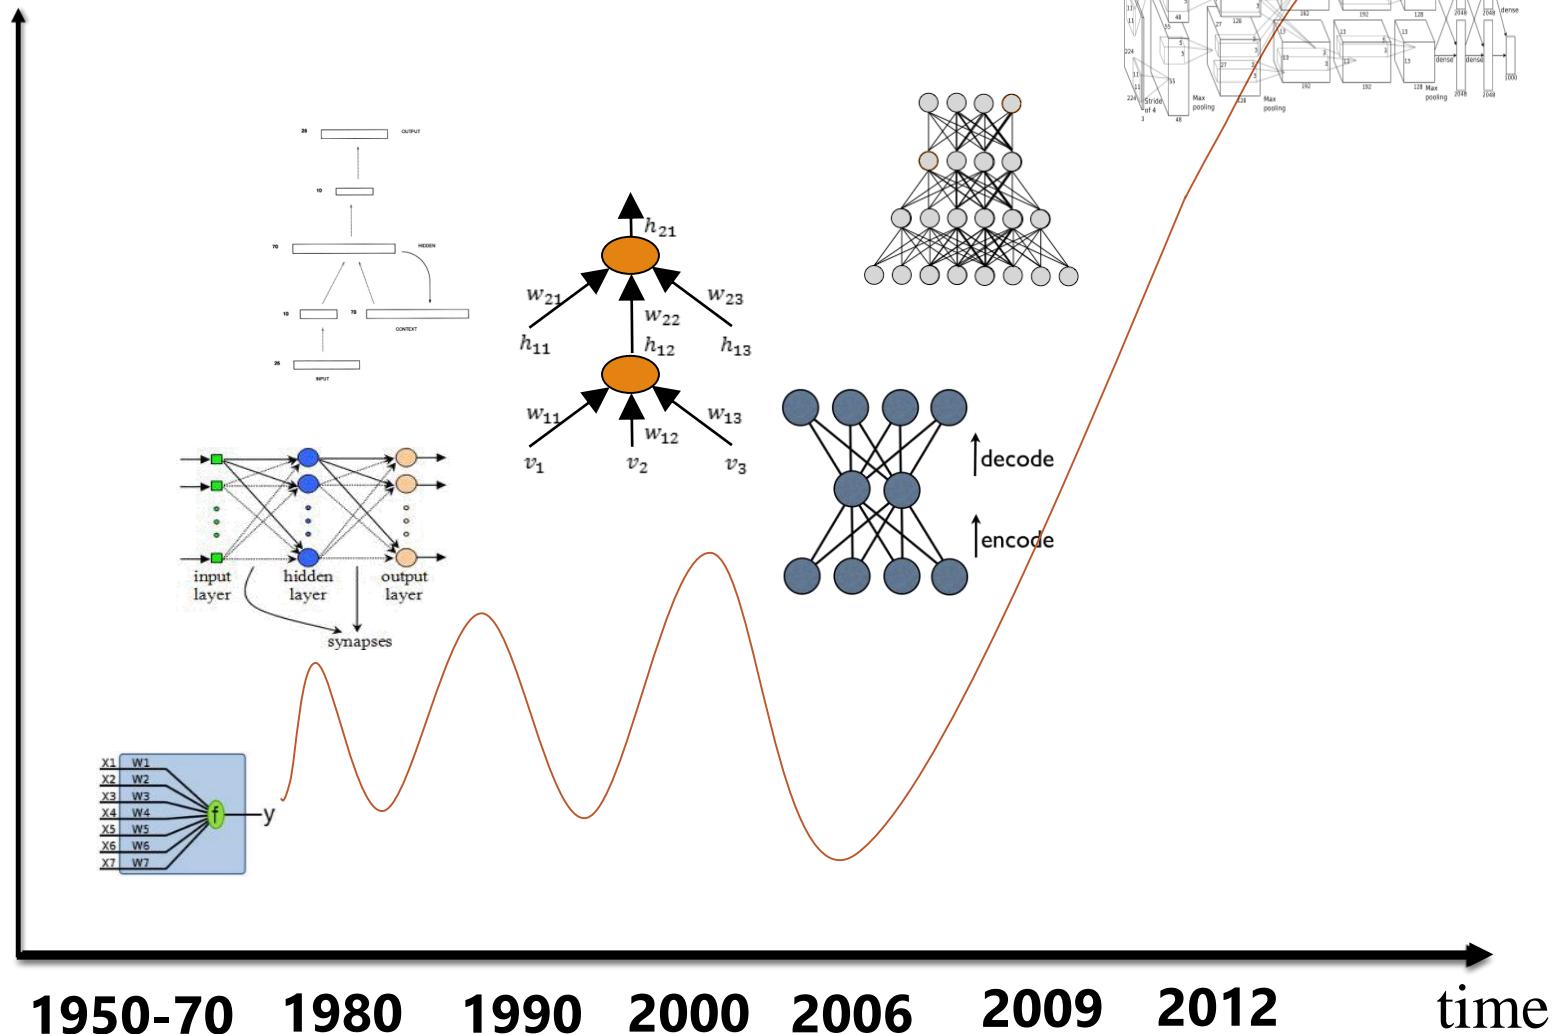
Hopfield J. (1982), "Neural networks and physical systems with emergent collective computational abilities", PNAS.

Rumelhart D., Hinton G. E., Williams R. J. (1986), "Learning internal representations by error-propagation".

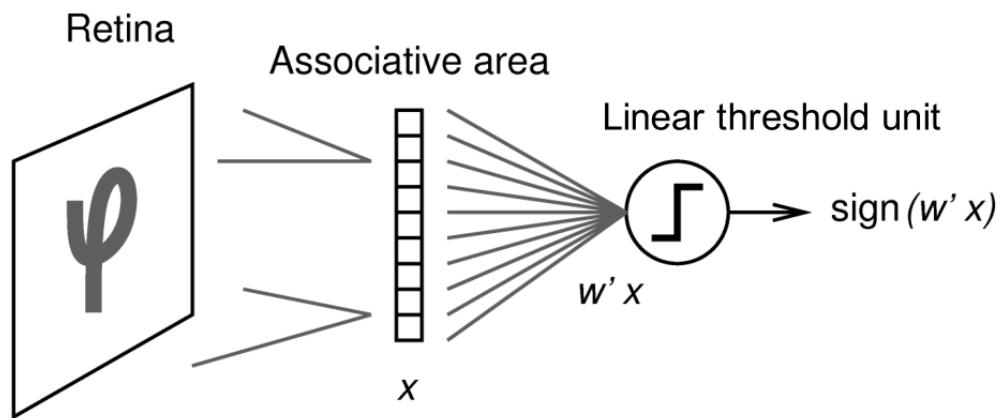
Elman, J.L. (1990). "Finding Structure in Time".

Hinton, G. E.; Osindero, S.; Teh, Y. (2006). "A fast learning algorithm for deep belief nets".

Krizhevsky, A.; Sutskever, I.; Hinton, G. (2012) "ImageNet Classification with Deep Convolutional Neural Networks"



Perceptron



Supervised learning of the weights w using the Perceptron algorithm.

The XOR problem that kills the Perceptron algorithm

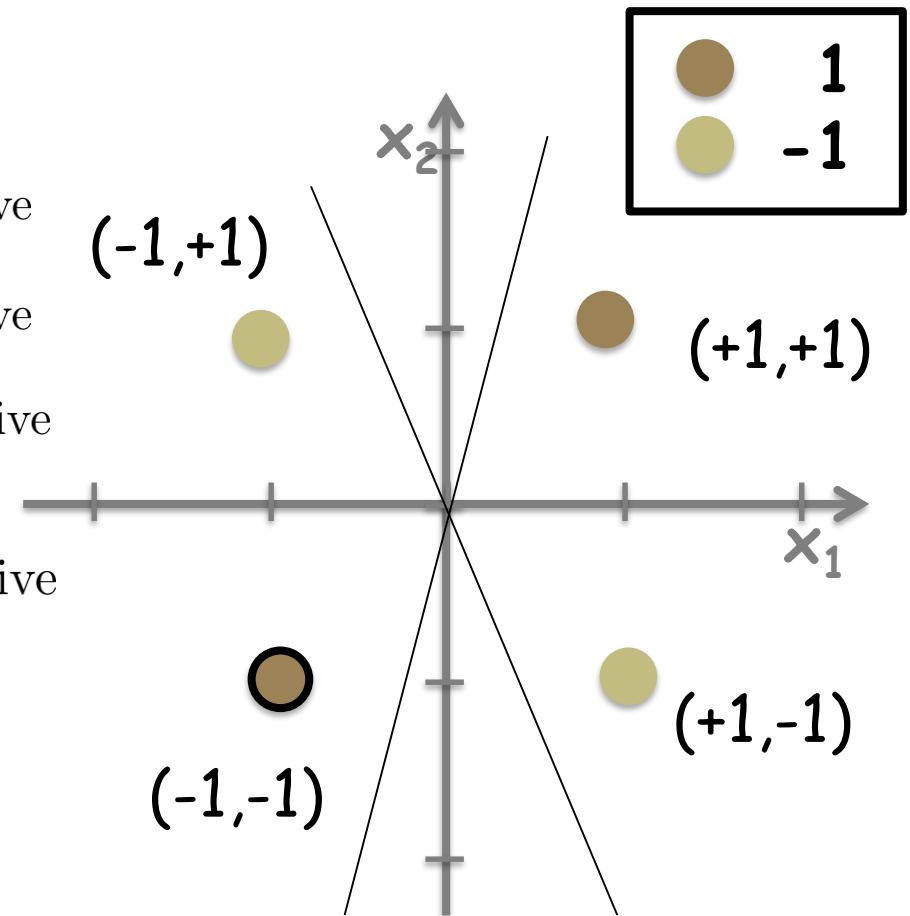
The XOR operator: \oplus

Positive (+1) \oplus Negative (-1) = Negative

Positive (+1) \oplus Positive (+1) = Positive

Negative (-1) \oplus Positive (+1) = Negative

Negative (-1) \oplus Negative (-1) = Positive



AI after Minsky's book *Perceptron*



Applications of deep learning technologies



10 BREAKTHROUGH TECHNOLOGIES 2013

Deep Learning

With massive amounts of computational power, machines can now recognize objects and translate speech in real time. Artificial intelligence is finally getting smart.



Welcome to Microsoft Project Oxford
An evolving portfolio of REST APIs and SDKs enabling developers to easily add intelligent services into their solutions to leverage the power of Microsoft's natural data understanding

Subscribe and get started for free today

Vision

- Computer Vision APIs BETA
Understand images and generate thumbnails
- Face APIs BETA
See your users with Face Detection and Recognition
- Emotion APIs BETA
Understand your users with Emotion Recognition

Speech

- Speech APIs BETA
Communicate with your users with speech recognition and synthesis powered by Bing

Language

- Spell Check APIs BETA
Deep and comprehensive uncommon spelling errors, via the Bing document index
- Language Understanding Intelligent Service (Luis) BETA
Understand natural language commands tailored to your application



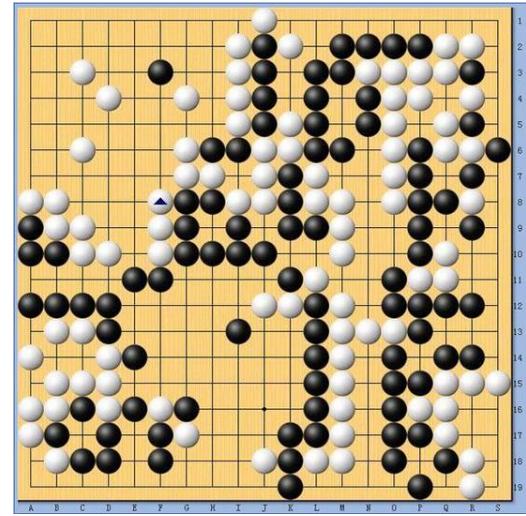
AlphaGo vs. Lee Sedol



AlphaGo



Lee Sedol



2016: match 1

AlphaGo: white; Lee Sedol: black



2017: AlphaGo vs. Jie Ke

Comparison to the 2016 version:

1. Trained (reinforcement learning) from scratch without human knowledge.
2. 10 times stronger than the 2016 version (can give up three moves).
3. Run on a single machine (10 TPUs).
4. Consumes 1/10 of the power.

Deep Learning

Geoffrey Hinton started this wave of development of deep learning in 2006.

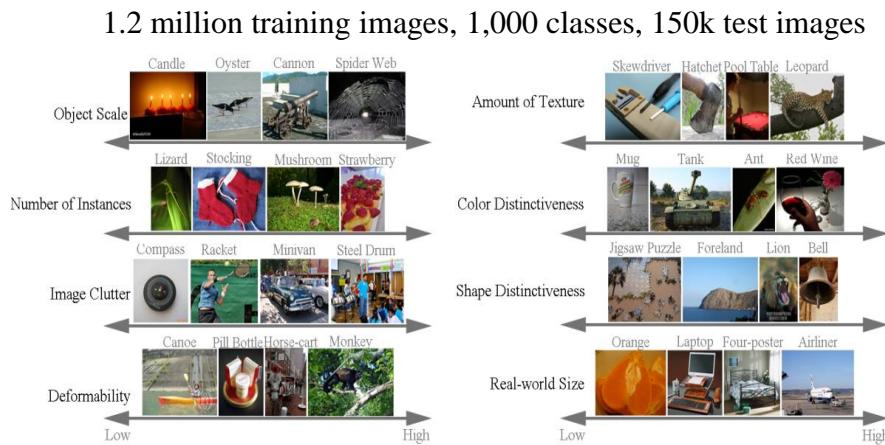
A large number of applications are mainly based on convolutional neural networks based on the models proposed by Yann LeCun.

Yoshua Bengio also made a big contribution in deep learning by proposing many fundamental algorithms.



Factors in the success of deep learning methods

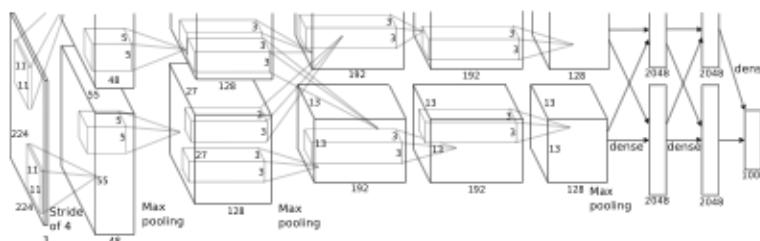
Availability of large amount of training data (big data), e.g. ImageNet.



Access of modern computing infrastructures, e.g. Nvidia GPUs



New development in neural networks with deep structures, e.g. AlexNet.





"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."



"a young boy is holding a baseball bat."

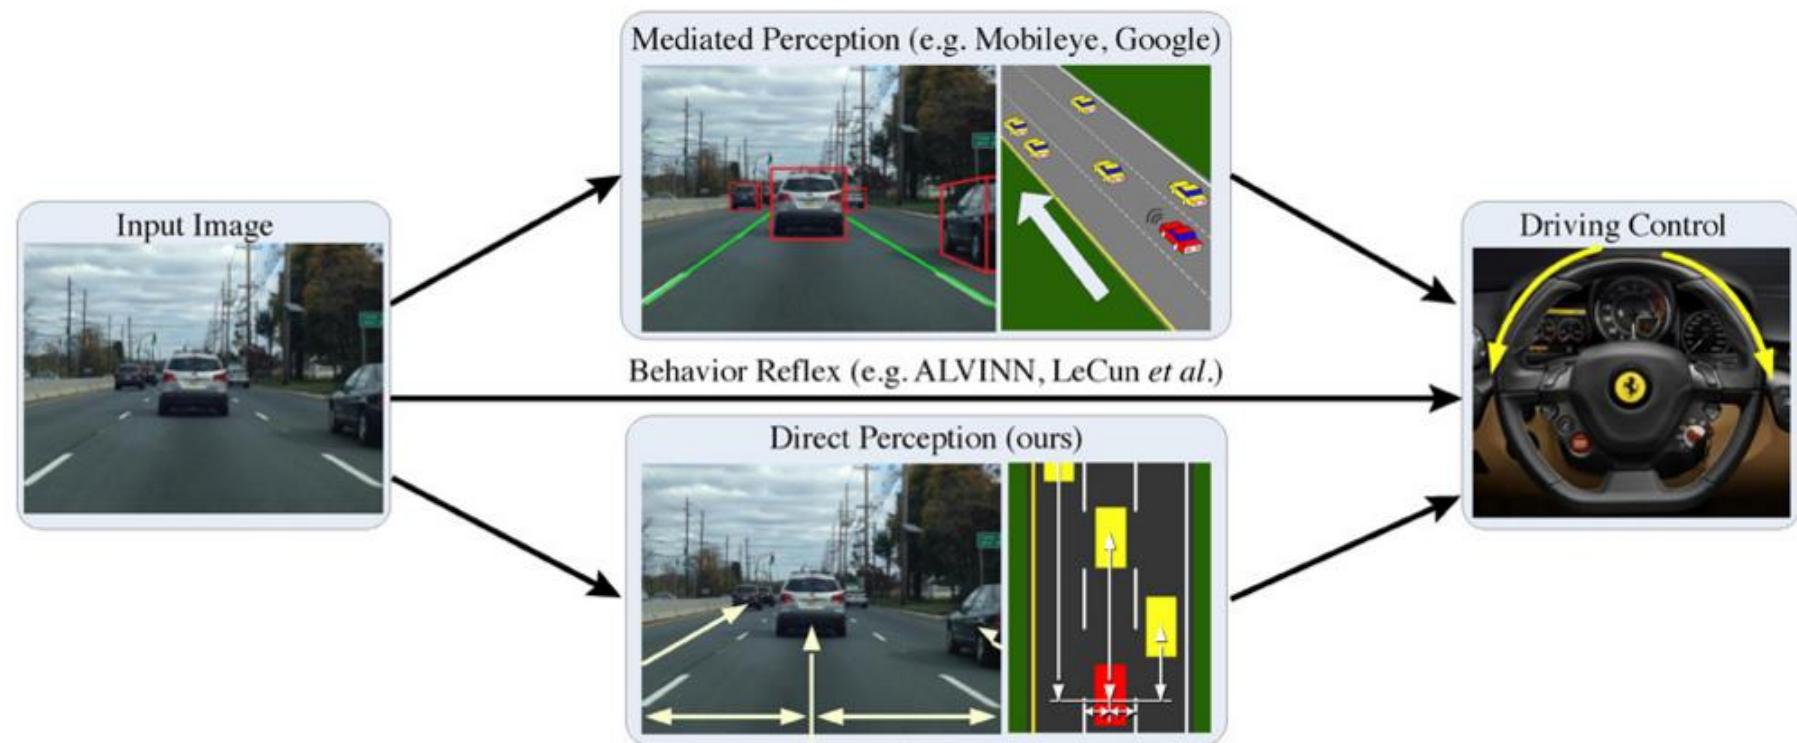


"a cat is sitt ing on a couch with
a remote controL"



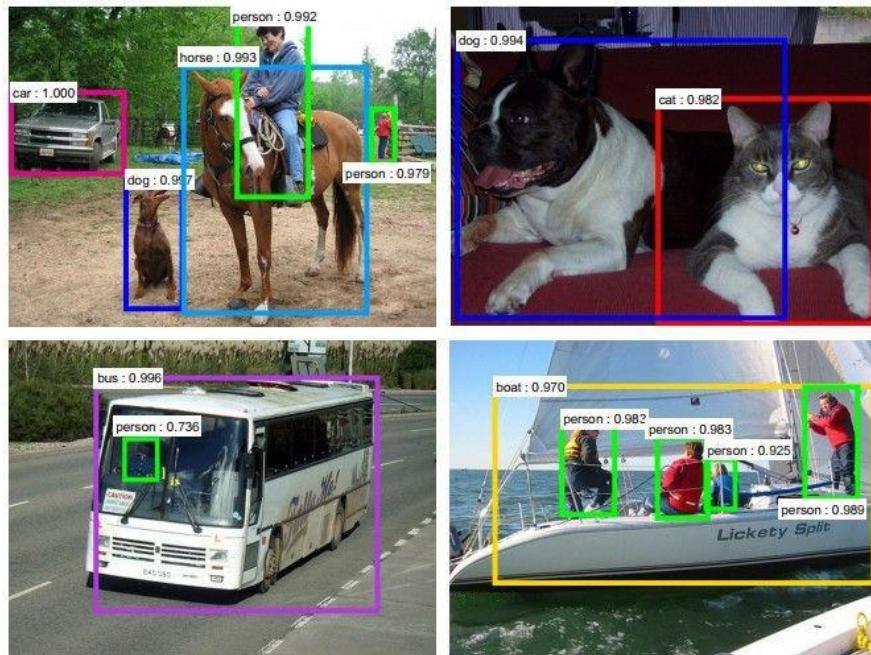
"a horse is standing in the middle of a
road."

DeepDriving: Learning Affordance for Direct Perception in Autonomous Driving



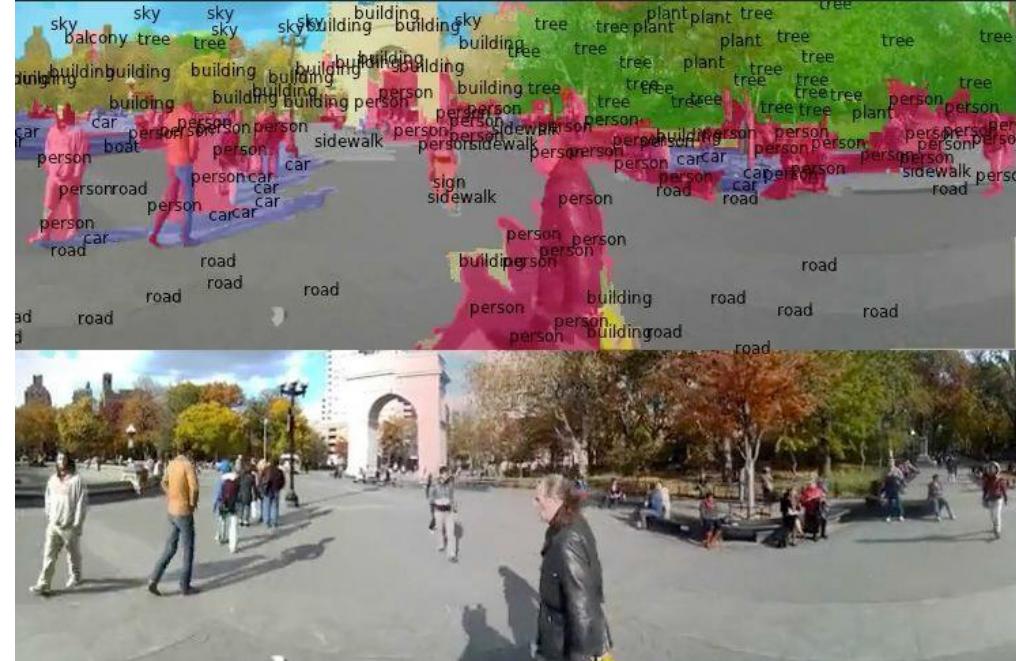
Fast-forward to today: ConvNets are everywhere

Detection



[Faster R-CNN: Ren, He, Girshick, Sun 2015]

Segmentation



[Farabet et al., 2012]

Neural Artistic Style Learning

Leon A. Gatys, Alexander S. Ecker, Matthias Bethge



Neural Audio Synthesis of Musical Notes with WaveNet Autoencoders, Engel et al. 2017

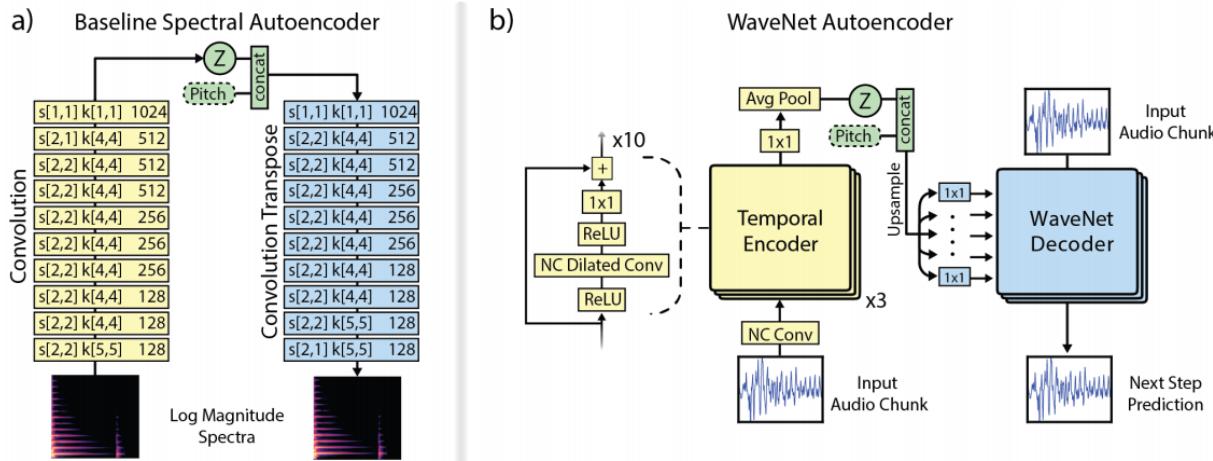
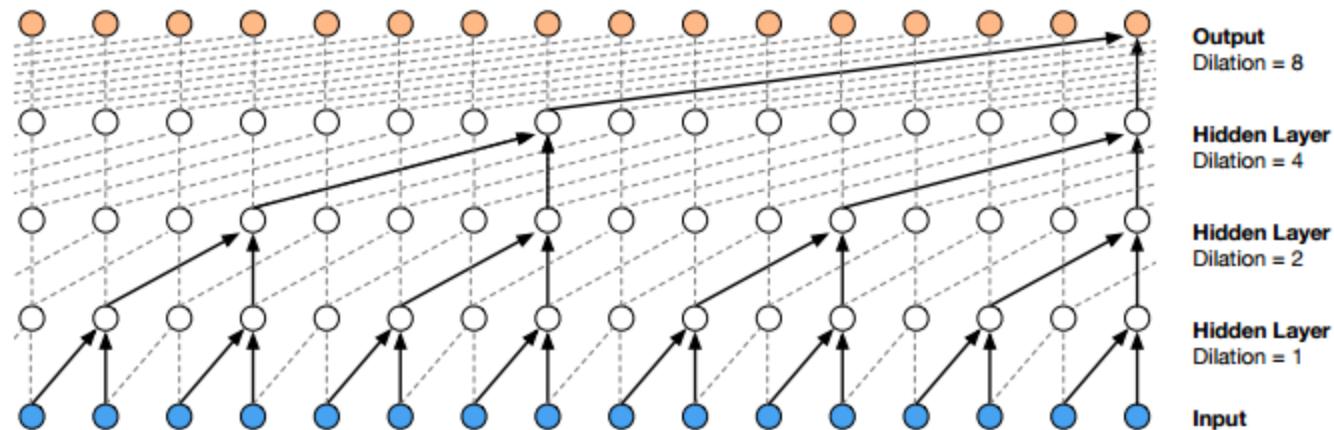


Figure 1. Models considered in this paper. For both models, we optionally condition on pitch by concatenating the hidden embedding with a one-hot pitch representation. 1a. Baseline spectral autoencoder: Each block represents a nonlinear 2-D convolution with stride (s), kernel size (k), and channels ($\#$). 1b. The WaveNet autoencoder: Downsampling in the encoder occurs only in the average pooling layer. The embeddings are distributed in time and upsampled with nearest neighbor interpolation to the original resolution before biasing each layer of the decoder. ‘NC’ indicates non-causal convolution. ‘ 1×1 ’ indicates a 1-D convolution with kernel size 1. See Section 2.1 for further details.



WAVENET: A GENERATIVE MODEL FOR RAW AUDIO, van den Oord et al. 2016

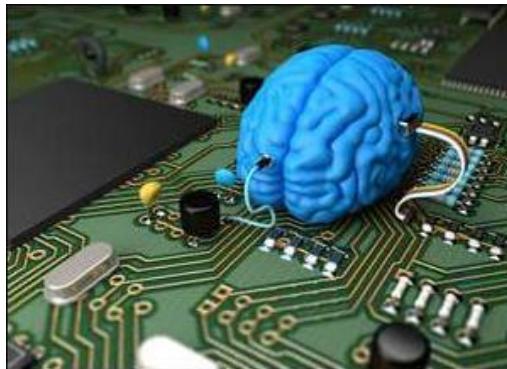


Sports Video Classification

From Yann LeCun



What society thinks I do



What my friends think I do



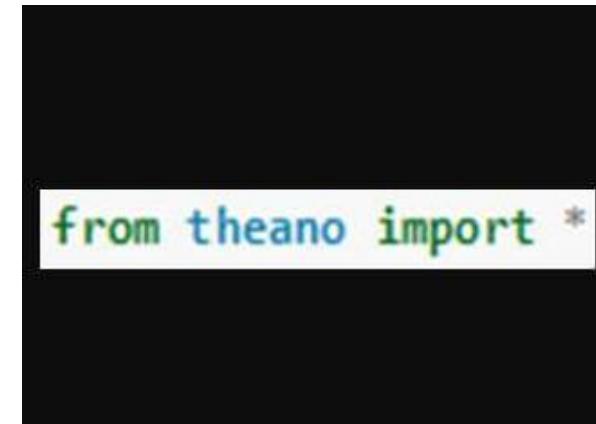
What other computer
scientists think I do



What mathematicians think I do



What I think I do



What I actually do

A bit of history:

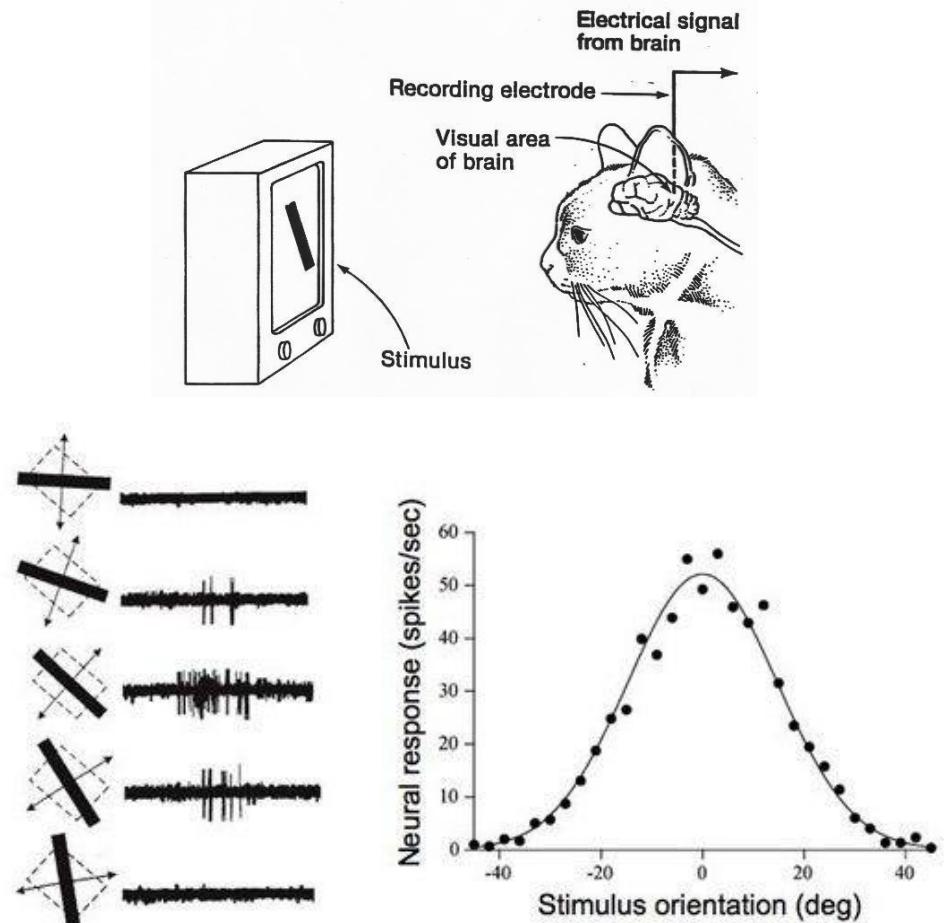
**Hubel & Wiesel,
1959**

RECEPTIVE FIELDS OF SINGLE
NEURONES IN
THE CAT'S STRIATE CORTEX

1962

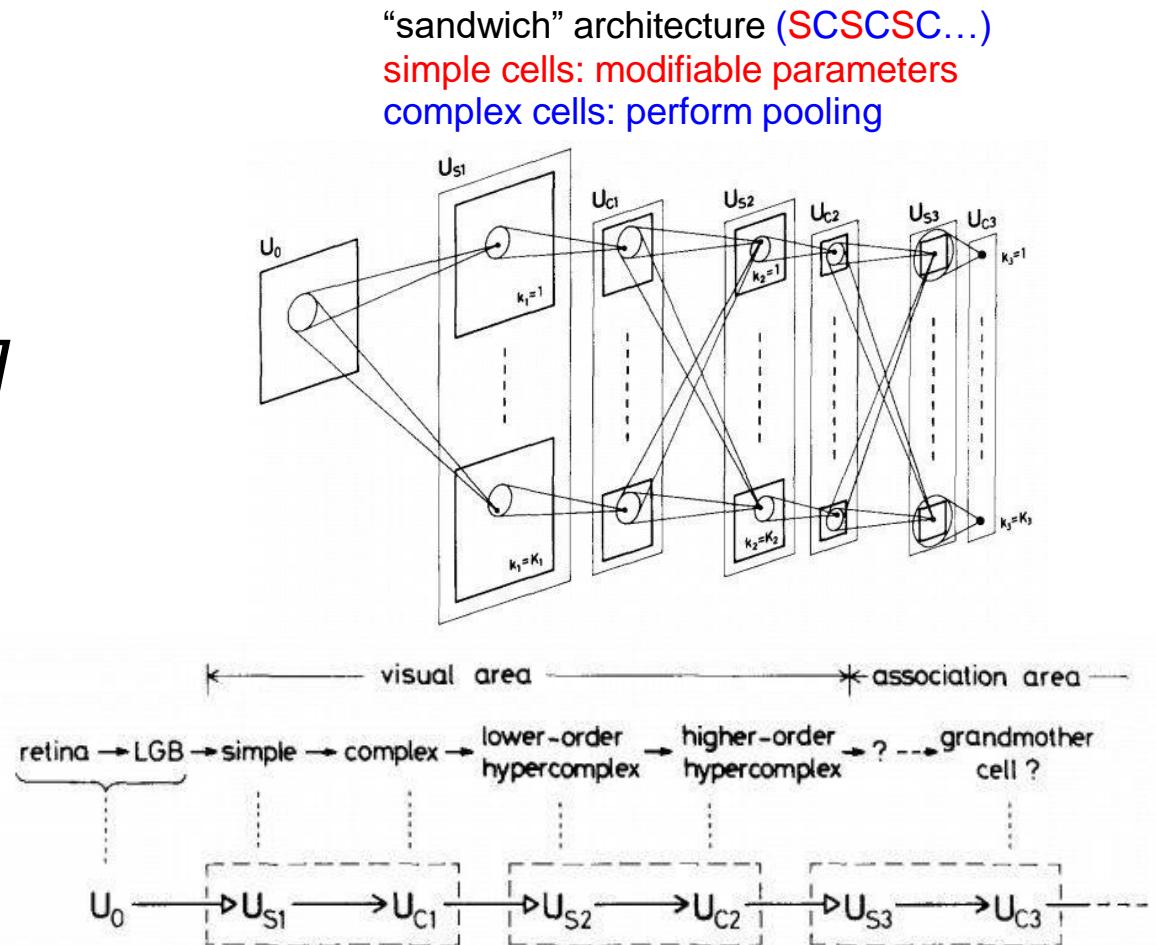
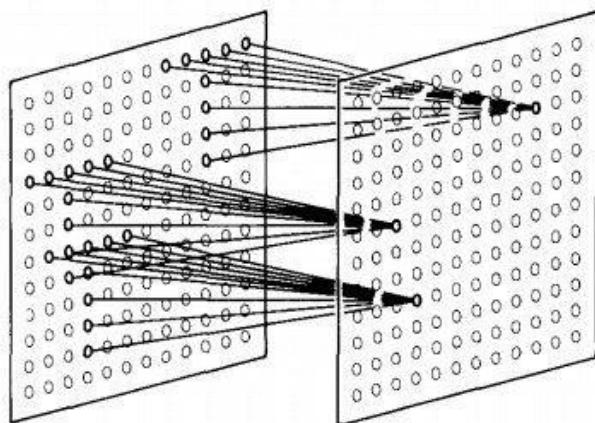
RECEPTIVE FIELDS, BINOCULAR
INTERACTION
AND FUNCTIONAL ARCHITECTURE IN
THE CAT'S VISUAL CORTEX

1968...



A bit of history:

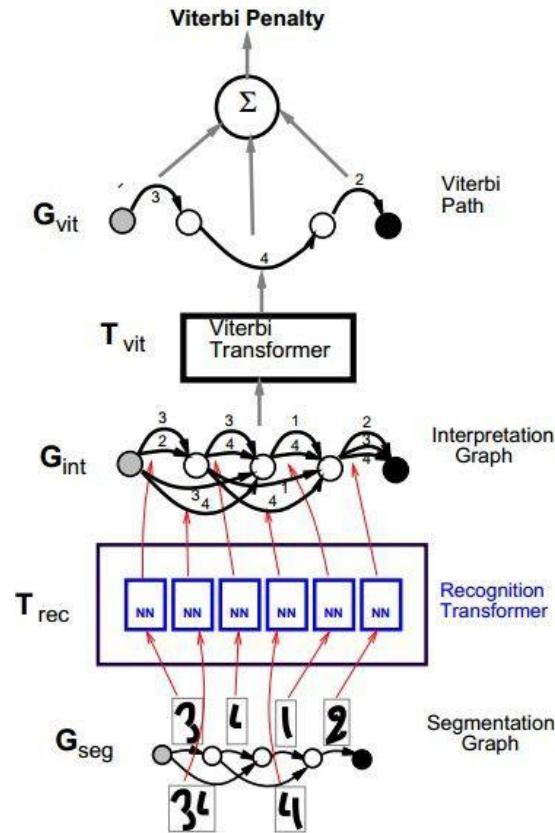
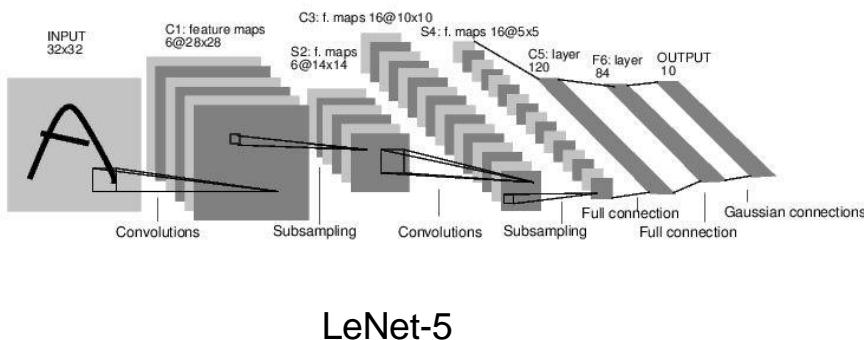
Neurocognitron [Fukushima 1980]



A bit of history:

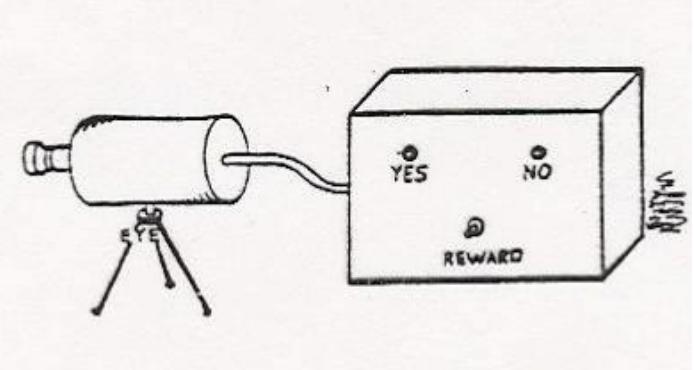
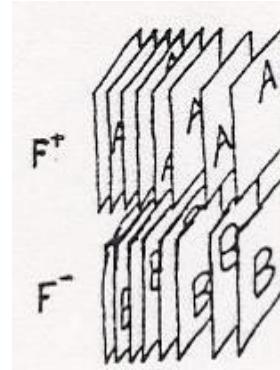
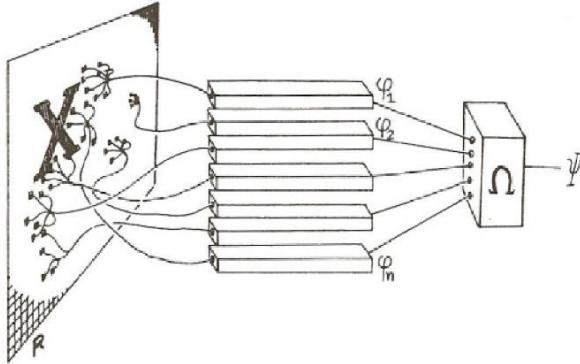
Gradient-based learning applied to document recognition

[LeCun, Bottou, Bengio, Haffner
1998]



Artificial neural networks: a brief history

- 1962
 - Frank Rosenblatt, *Principles of Neurodynamics: Perceptrons and the Theory of Brain Mechanisms*
 - Perceptron can learn anything you can program it to do.

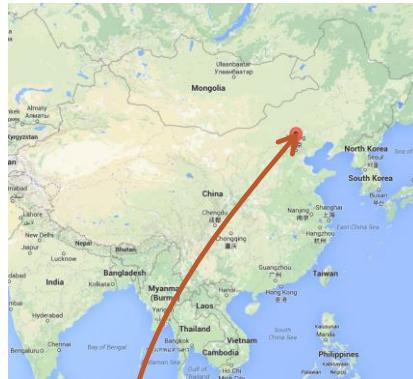


1969

Minsky & Papert, Perceptrons: An introduction to computational geometry

There are many things a perceptron can't in principle learn to do

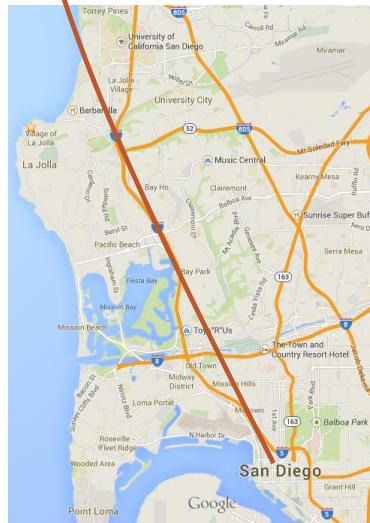
An illustration



Beijing, China

Our task: to learn from a data (but not always)

representation: vehicles (invention) +
function: itinerary (composition)



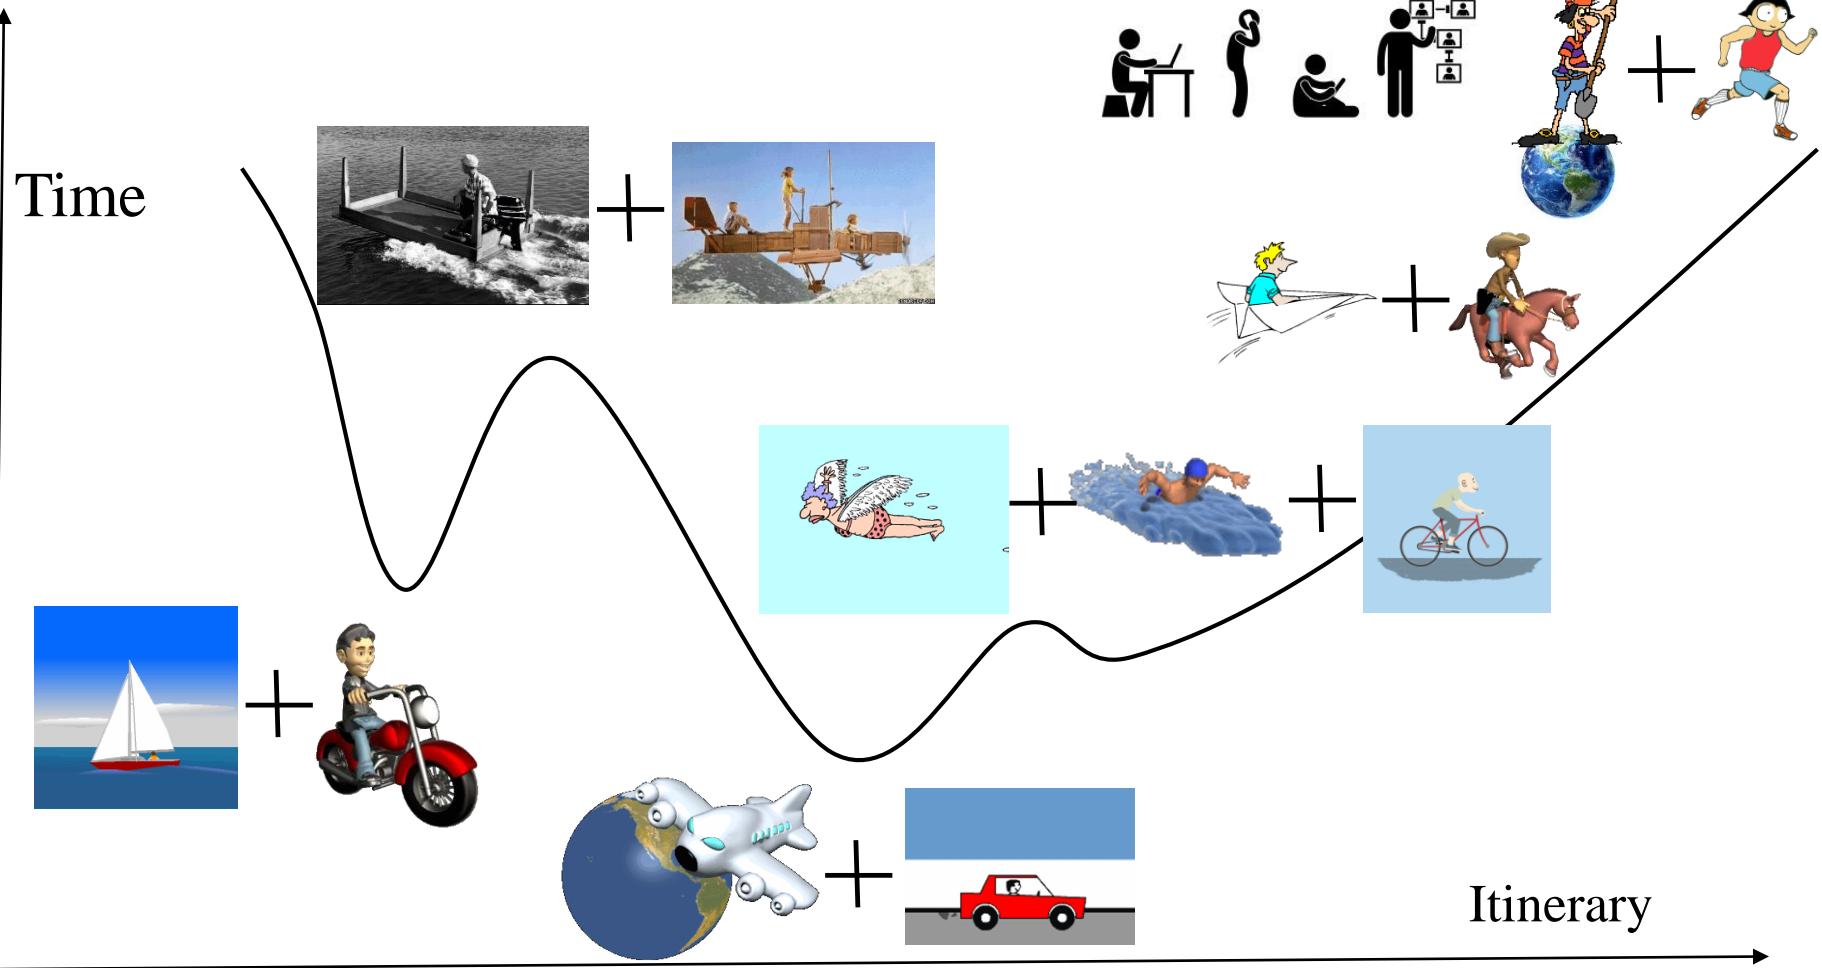
San Diego, USA

An illustration

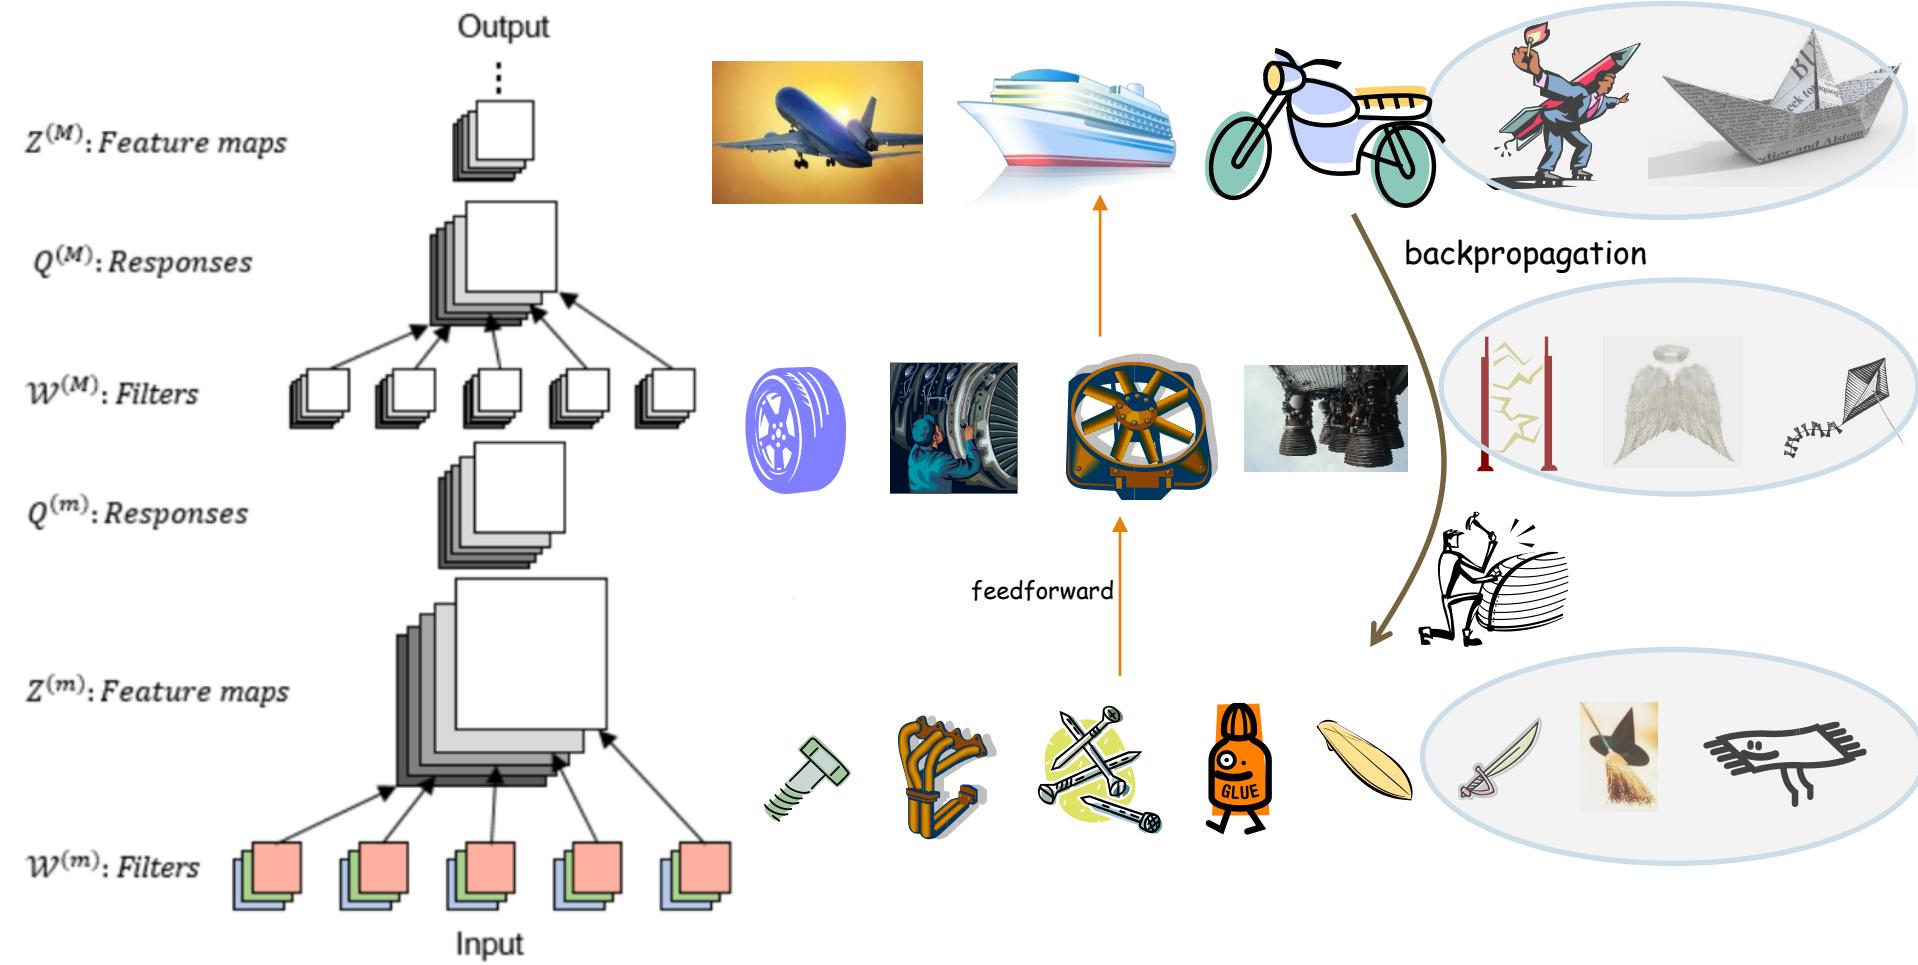
(images obtained by searching Google and Bing)

Purpose of a data-driven learning approach:

- (1) Inventing--features
- (2) Composing--features

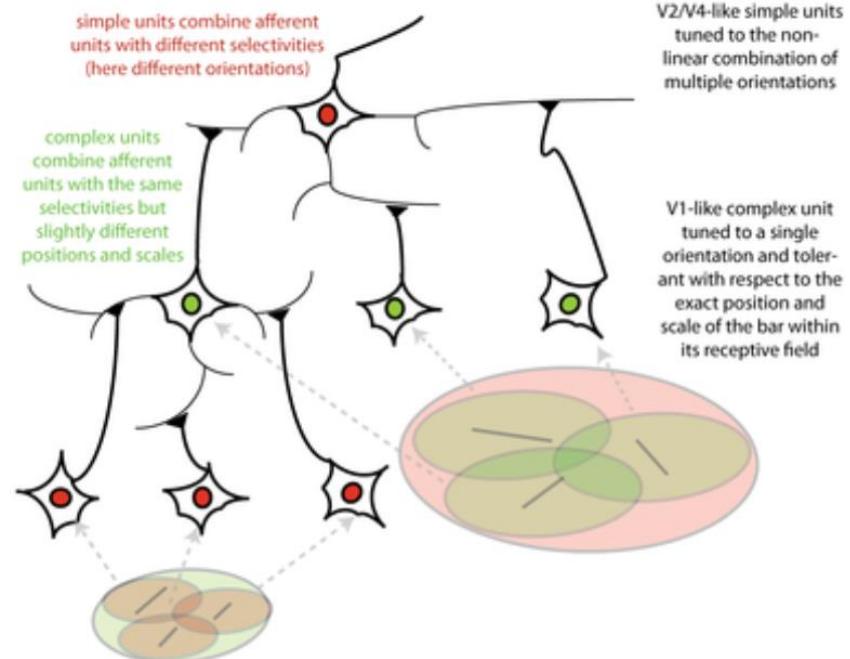
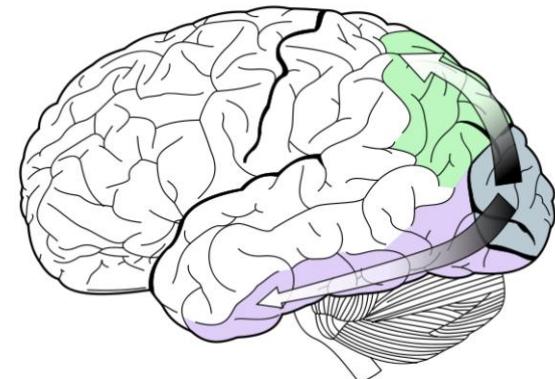


An illustration



Convolutional neural networks (Y. LeCun)

Visual Representation

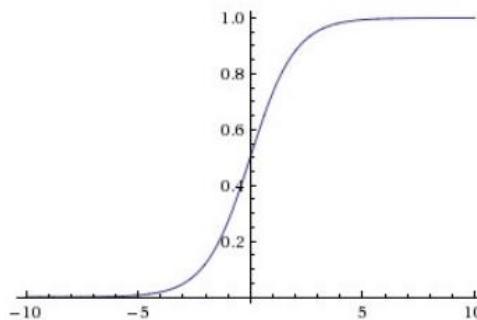
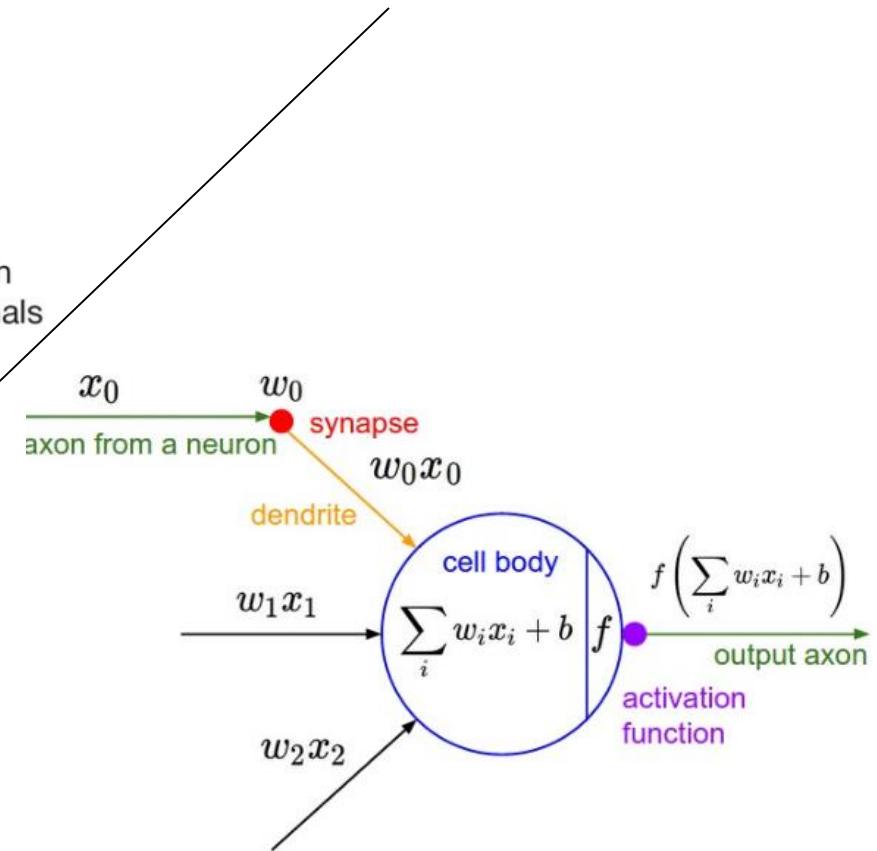
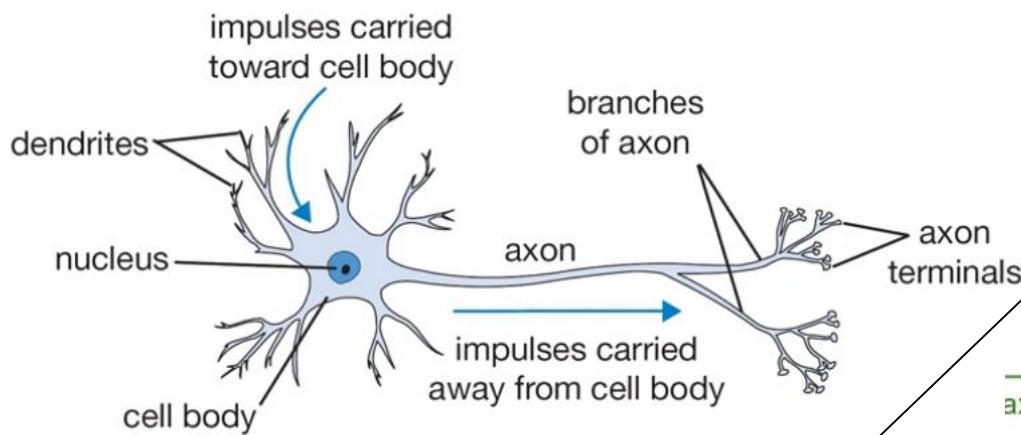


V2	V4	posterior IT	anterior IT

Hubel and Wiesel Model

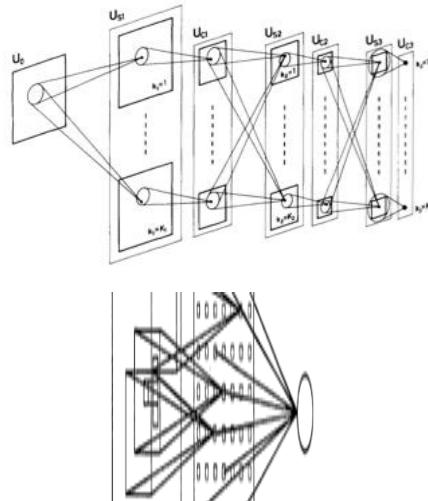
Kobatake and Tanaka, 1994

Perceptron



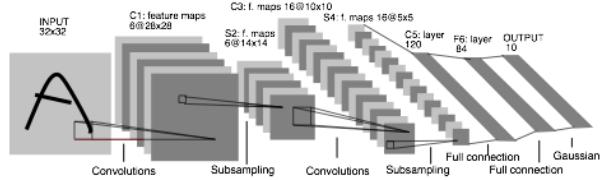
Sigmoid function
$$f(x) = \frac{1}{1+e^{-x}}$$

History of ConvNets

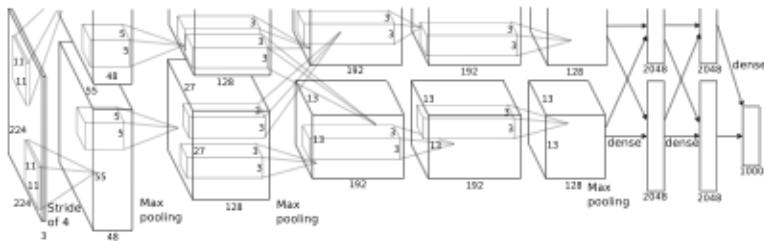


Fukushima 1980
Neocognitron

Rumelhart, Hinton, Williams 1986
“T” versus “C” problem



LeCun et al. 1989-1998
Hand-written digit reading



Krizhevksy, Sutskever, Hinton 2012
ImageNet classification

ImageNet experiments

